



# Intro to Visualization in Python - Static Plots - 1

*One should look for what is and not what he thinks should be. (Albert Einstein)*

# Module completion checklist

Objective	Complete
Prepare data for visualization	
Create histograms, boxplots, and bar charts	

# Visualizing data with matplotlib



- `matplotlib` is a popular plotting library among scientists and data analysts
- It is one of the older Python plotting libraries, so it has become quite flexible and **well-documented**
- Other plotting libraries you may come across are Seaborn (which is built on `matplotlib`), ggplot (the Python version of the popular R plotting library), Plotly, Bokeh, and many others
- Pandas also come with some plotting capabilities, and these are just based on `matplotlib`
- Explore the different types of plots you can create with `matplotlib` by browsing their **gallery**

# Loading packages

- Let's load the packages we will be using:

```
import pandas as pd
import numpy as np
import pickle
import os
from pathlib import Path
```

# Directory settings

- In order to maximize the efficiency of your workflow, you should encode your directory structure into variables
- We will use the `pathlib` library
- Let the `main_dir` be the variable corresponding to your course folder
- Let `data_dir` be the variable corresponding to your data folder

```
# Set 'main_dir' to location of the project folder
home_dir = Path(".").resolve()
main_dir = home_dir.parent.parent
print(main_dir)
```

```
data_dir = str(main_dir) + "/data"
print(data_dir)
```

```
plot_dir = str(main_dir) + "/plots"
if not os.path.exists(plot_dir):
    os.makedirs(plot_dir)
print(plot_dir)
```

# Importing matplotlib

- Let's import `pyplot` as `plt` so that we can call `plt.[any_function]()` with appropriate arguments to create a plot
- The `pyplot` module of the `matplotlib` library has a large and diverse set of functions
- It allows us to create pretty much any conceivable visualization out there
- See the documentation on `pyplot` [here](#)

```
import matplotlib.pyplot as plt
```

## matplotlib.pyplot

`matplotlib.pyplot` is a state-based interface to matplotlib. It provides a MATLAB-like way of plotting.

`pyplot` is mainly intended for interactive plots and simple cases of programmatic plot generation:

```
import numpy as np
import matplotlib.pyplot as plt

x = np.arange(0, 5, 0.1)
y = np.sin(x)
plt.plot(x, y)
```

The object-oriented API is recommended for more complex plots.

## Functions

<code>acorr(x, *[, data])</code>	Plot the autocorrelation of <code>x</code> .
<code>angle_spectrum(x[, Fs, Fc, window, pad_to, ...])</code>	Plot the angle spectrum.
<code>annotate(s, xy, *args, **kwargs)</code>	Annotate the point <code>xy</code> with text <code>s</code> .
<code>arrow(x, y, dx, dy, **kwargs)</code>	Add an arrow to the axes.
<code>autoscale([enable, axis, tight])</code>	Autoscale the axis view to the data (toggle).

# Dataset for visualization

- We will now load the dataset and save it as df

```
# This dataset is of type dataframe. Let's assign this dataset to a variable, so that we can  
manipulate it freely.  
df = pd.read_csv(str(data_dir)+"/" + "diabetes.csv")
```

```
print(type(df))  #<- a Pandas DataFrame!
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
print(len(df))  #<- returns the number of rows
```

```
768
```

# Subsetting data

- Let's subset our data so that we have the variables we need
- Let's name this subset `df_subset`

```
df_subset = df[['BMI', 'SkinThickness', 'BloodPressure', 'Age', 'Glucose', 'Insulin',  
'DiabetesPedigreeFunction', 'Outcome', 'Pregnancies']]  
print(df_subset.head())
```

	BMI	SkinThickness	...	Outcome	Pregnancies
0	33.6	35	...	1	6
1	26.6	29	...	0	1
2	23.3	0	...	1	8
3	28.1	23	...	0	1
4	43.1	35	...	1	0

[5 rows x 9 columns]

- We are choosing these variables because they illustrate the concepts best
- However, you should be able to work with (and visualize) all of your data



# Data Reshaping: wide vs. long

- Talking about data reshaping usually refers to converting between what is called either **wide** or **long** data formats
  - **Wide** data is much more visually digestible, which is why you're likely to come across it if you are using data from some type of report
  - **Long** data is much easier to work with in Pandas, and generally speaking in most data analysis and plotting tools

# Data reshaping: wide vs long (cont'd)

- **Wide data** often appears when the values are some type of aggregate (we will use the mean of groups)
- Let's make a typical **wide dataframe** of two rows and six columns that looks like this:

```
Outcome      BMI      ...      Insulin  DiabetesPedigreeFunction
0          0  30.304200  ...    68.792000             0.429734
1          1  35.142537  ...   100.335821             0.550500

[2 rows x 8 columns]
```

# Prepare data: group and summarize

- Now that we know how to group and summarize data, let's create a summary dataset that would include the following:
  - Grouped data by `Target` variable
  - Mean value computed on the grouped data that includes the following variables:
    - `BMI`
    - `SkinThickness`
    - `BloodPressure`
    - `Age`
    - `Glucose`
    - `Insulin`
    - `DiabetesPedigreeFunction`

# Prepare data: group and summarize (cont'd)

- For demonstration, we use the original dataframe `df` to identify the grouping column
- Then we use this column to perform the groupby operation and find the mean of the columns present in `df_subset`

```
col_dict = df_subset.nunique().to_dict()
grouping_col = min(col_dict, key=col_dict.get)
# Group data by variable with min levels.
grouped = df_subset.groupby(grouping_col)
```

```
# Compute mean on the listed variables using the grouped data.
df_grouped_mean = grouped.mean()[['BMI', 'SkinThickness', 'BloodPressure', 'Age', 'Glucose',
'Insulin', 'DiabetesPedigreeFunction']]
print(df_grouped_mean)
```

	BMI	SkinThickness	...	Insulin	DiabetesPedigreeFunction
Outcome			...		
0	30.304200	19.664000	...	68.792000	0.429734
1	35.142537	22.164179	...	100.335821	0.550500

[2 rows x 7 columns]

# Prepare data: group and summarize (cont'd)

```
# Reset index of the dataset.  
df_grouped_mean = df_grouped_mean.reset_index()  
print(df_grouped_mean)
```

```
Outcome      BMI      ...      Insulin  DiabetesPedigreeFunction  
0          0  30.304200  ...    68.792000             0.429734  
1          1  35.142537  ...   100.335821             0.550500  
  
[2 rows x 8 columns]
```

- We call this dataframe **wide** because each variable has its own column
- It makes the table easier to present, but inconvenient to run analyses on or visualize


# Why long?

- Now let's convert this wide data to the **long format**
  - We are going to leave the `categorical` variable and the mean values as is in their columns
  - All of our other variables will appear as a single metric column
- This format is convenient to work with when we run analysis and plot the data

	Outcome	metric	mean
0	0	BMI	30.304200
1	1	BMI	35.142537
2	0	SkinThickness	19.664000
3	1	SkinThickness	22.164179
4	0	BloodPressure	68.184000

# Wide to long format: melt

- To **convert from wide to long format**, we use the Pandas `melt` function with the following arguments:
  - i. Wide dataframe
  - ii. Variable(s) that will be preserved as the `ids` of the data (like categorical variables)
  - iii. Name of the variable that will now contain the column names from the wide data we want to melt together
  - iv. Name of the column that will contain respective values corresponding to the melted columns



# Wide to long format: melt (cont'd)

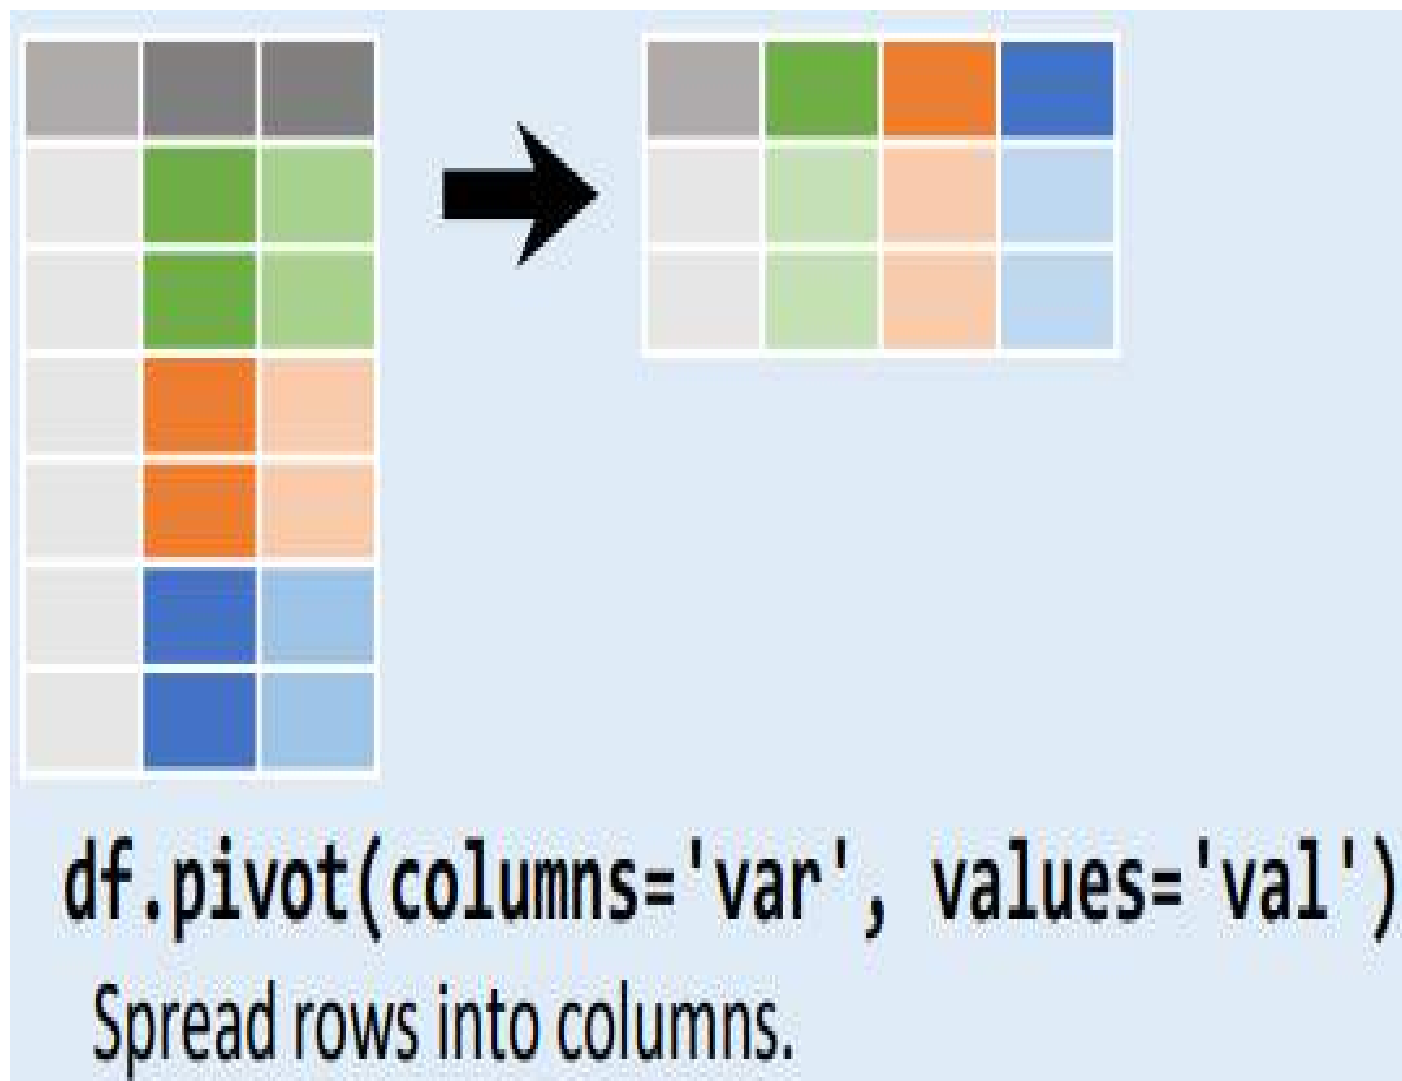
```
# Melt the wide data into long.
df_grouped_mean_long = pd.melt(df_grouped_mean,          #<- wide dataset
                               id_vars = [grouping_col],  #<- identifying variable
                               var_name = 'metric',       #<- contains col names of wide data
                               value_name = 'mean')       #<- contains values from above columns

print(df_grouped_mean_long)
```

	Outcome		metric	mean
0	0		BMI	30.304200
1	1		BMI	35.142537
2	0		SkinThickness	19.664000
3	1		SkinThickness	22.164179
4	0		BloodPressure	68.184000
5	1		BloodPressure	70.824627
6	0		Age	31.190000
7	1		Age	37.067164
8	0		Glucose	109.980000
9	1		Glucose	141.257463
10	0		Insulin	68.792000
11	1		Insulin	100.335821
12	0	DiabetesPedigreeFunction		0.429734
13	1	DiabetesPedigreeFunction		0.550500



# Long to wide format: pivot



We can convert the **long data back to wide** format with the `.pivot()` method

1. The `index` argument refers to what values will become the `ids` in the new dataframe
2. The `columns` argument refers to the column in which its values will be converted to column names
3. Lastly, we supply the `values` argument to fill in the values of the wide data

# Long to wide format: pivot (cont'd)

```
# Melt the long data into wide.
df_grouped_mean_wide = df_grouped_mean_long.pivot(
    index = [grouping_col],    #<- identifying
    variable = ...,           #<- variable names of long data
    columns = 'metric',       #<- col names of wide data
    values = 'mean')          #<- values from above
columns
print(df_grouped_mean_wide)
```

metric	Age	BMI	...	Insulin	SkinThickness
Outcome			...		
0	31.190000	30.304200	...	68.792000	19.664000
1	37.067164	35.142537	...	100.335821	22.164179

[2 rows x 7 columns]

# Module completion checklist

Objective	Complete
Prepare data for visualization	✓
Create histograms, boxplots, and bar charts	


# Univariate plots

- Univariate plots are used to **visualize the distribution of a single variable**
- They are mainly used in the initial stages of EDA when we want to learn more about individual variables in our data
- They are also combined with other univariate plots to compare data distributions of different variables
- Univariate plots include the following popular graphs: `histogram`, `boxplot`, `density curve`, `dot plot`, `QQ plot`, **and** `bar plot`

# Univariate plots: histogram

- A histogram represents the **distribution of numerical data**
- The height of each bar has been calculated as the number of observations in that range
- We can use `plt.hist()` to produce a basic histogram of any numeric variable


```
plt.rcParams.update({'font.size': 15})  
plt.hist(df_subset['BMI'])  
plt.show()
```



# Univariate plots: histogram (cont'd)

- Bins represent the intervals in which we want to group the observations
- Control the number of bins with the `bins` parameter
- As the **number of bins increases**, the range of values each bin represents decreases, and so does the height of the bar

```
plt.hist(df_subset['BMI'], bins = 20)
plt.xlabel('BMI')           #<- label x-axis
plt.ylabel('Frequency')     #<- label y-axis
plt.title('BMI distribution') #<- add plot title
plt.show()
```




# Univariate plots: boxplot

- A boxplot is a visual summary of the **25th, 50th, and 75th percentiles**
- The orange line shows the median of `ppl_total`
- The top and bottom of the box are the 25th and 75th percentile respectively
- The outermost lines are called the `whiskers`
- Values beyond whiskers are considered **outliers** - they are substantially outside the rest of the data

```
plt.boxplot(df_subset['BMI'])
```

```
plt.show()
```



# Univariate plots: boxplot (cont'd)

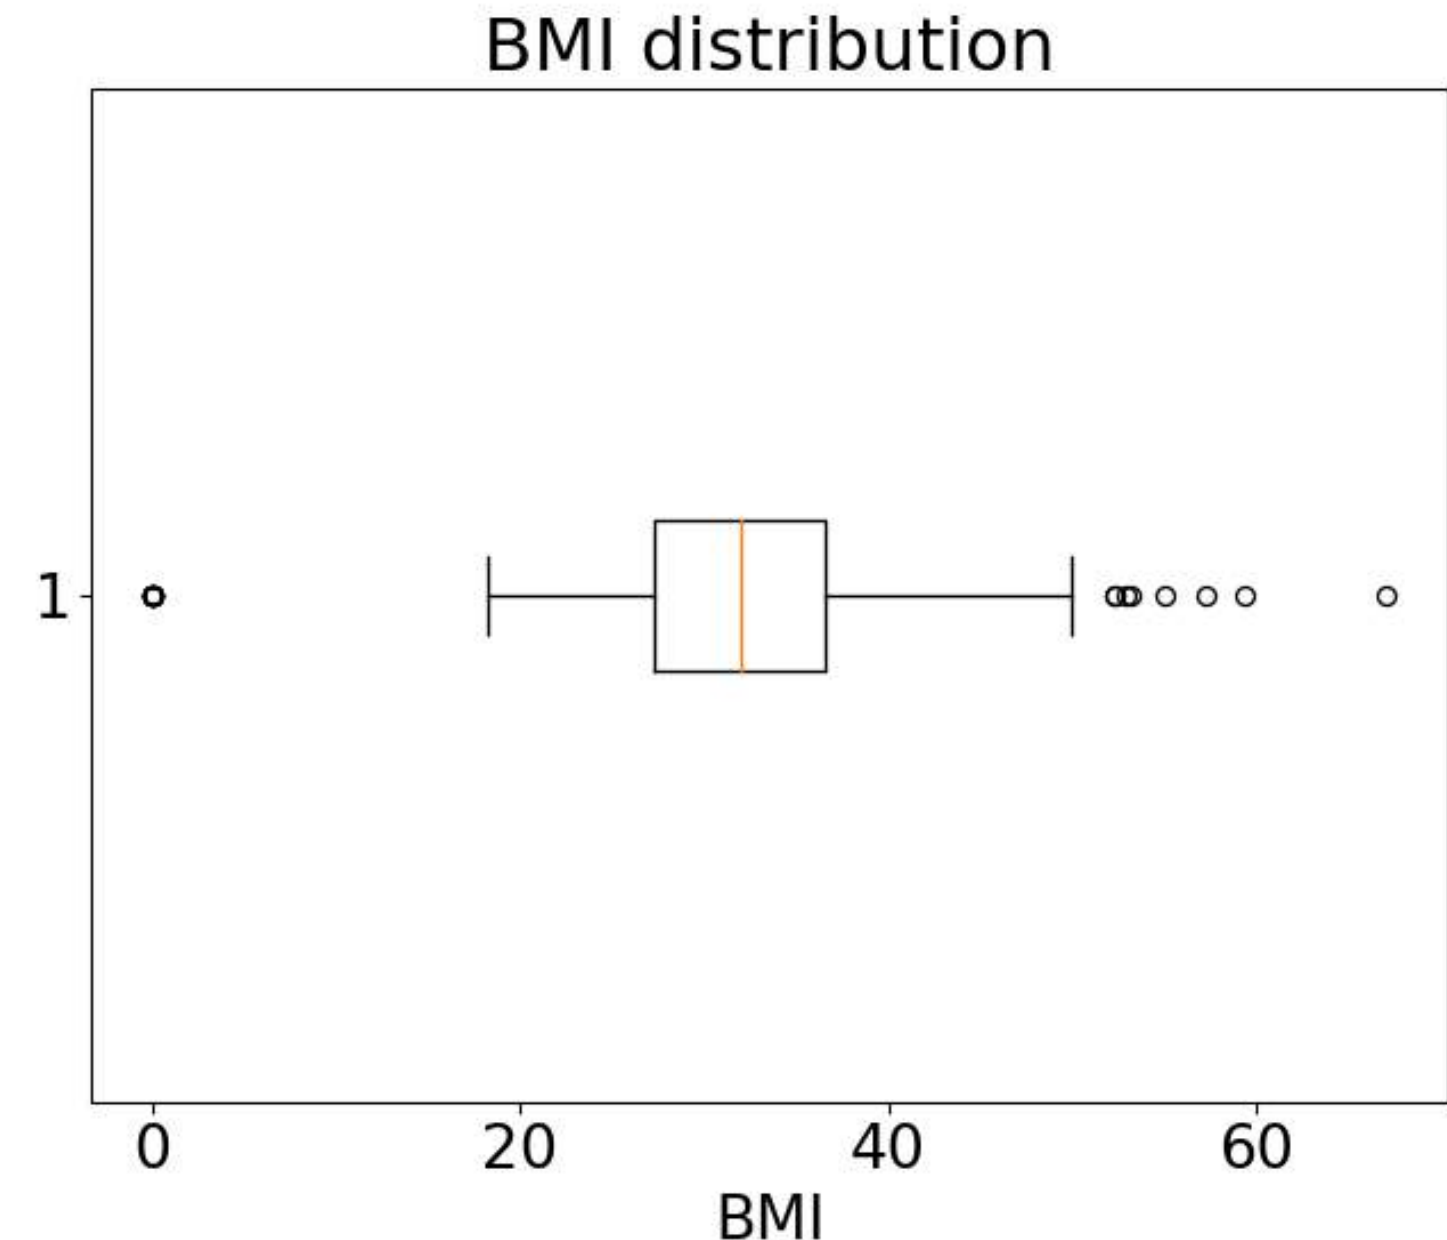
- You can change the orientation of the plot to horizontal by setting `vert = False`
- Answer in chat: By looking at this boxplot, what can you tell about the '**BMI**' distribution in our data?

```
plt.show()
```

```
plt.boxplot(df_subset['BMI'], vert = False)
```

```
plt.xlabel('BMI')           # label x-axis
```

```
plt.title('BMI distribution') # add plot title
```





# Univariate plots: bar chart

- A bar chart is a plot where the height of each bar represents **the numeric value of a category**
- We can use `plt.bar()` to produce a basic histogram of **any categorical variable**
- Bar charts are most commonly used when visualizing survey data or summary data
- The general syntax for creating a bar chart consists of 3 main variables:
  - position of the bars on the `axis`
  - height of the bars
  - names of categories that are used to label the bars

# Univariate plots: bar chart - cont'd

- When plotting bar charts of any complexity, the best type of data to use is **long data**
- First create a simple bar chart of the variable means using the `df_grouped_mean_long` data we created earlier

```
print(df_grouped_mean_long.head())
```

	Outcome	metric	mean
0	0	BMI	30.304200
1	1	BMI	35.142537
2	0	SkinThickness	19.664000
3	1	SkinThickness	22.164179
4	0	BloodPressure	68.184000

# Univariate plots: bar chart - cont'd

- Next, filter 'Outcome' by a category and only keep two columns: `metric` and `mean`

```
query = 'Outcome' + "==" + str('0')
df_true_means = df_grouped_mean_long.query(query) [['metric', 'mean']]
print(df_true_means)
```

	metric	mean
0	BMI	30.304200
2	SkinThickness	19.664000
4	BloodPressure	68.184000
6	Age	31.190000
8	Glucose	109.980000
10	Insulin	68.792000
12	DiabetesPedigreeFunction	0.429734

# Univariate plots: bar chart - cont'd

- Now, get the data we need and assign it to the three variables for convenience and clarity:
1. **Categories** (i.e., labels) that will represent each bar are all contained in the `metric` column
  2. **Bar heights** are contained in the `mean` column for each of the 5 categories
  3. **Bar positions** will be a range of numbers based on the number of categories (i.e., bars)

```
bar_labels = df_true_means['metric']      #<- 1
bar_heights = df_true_means['mean']       #<- 2
num_bars = len(bar_heights)
bar_positions = np.arange(num_bars)       #<- 3
```

# Univariate plots: bar chart - cont'd

- Labels are tricky to fit sometimes, so we can either **adjust** the figure size or label orientation

```
plt.show()
```

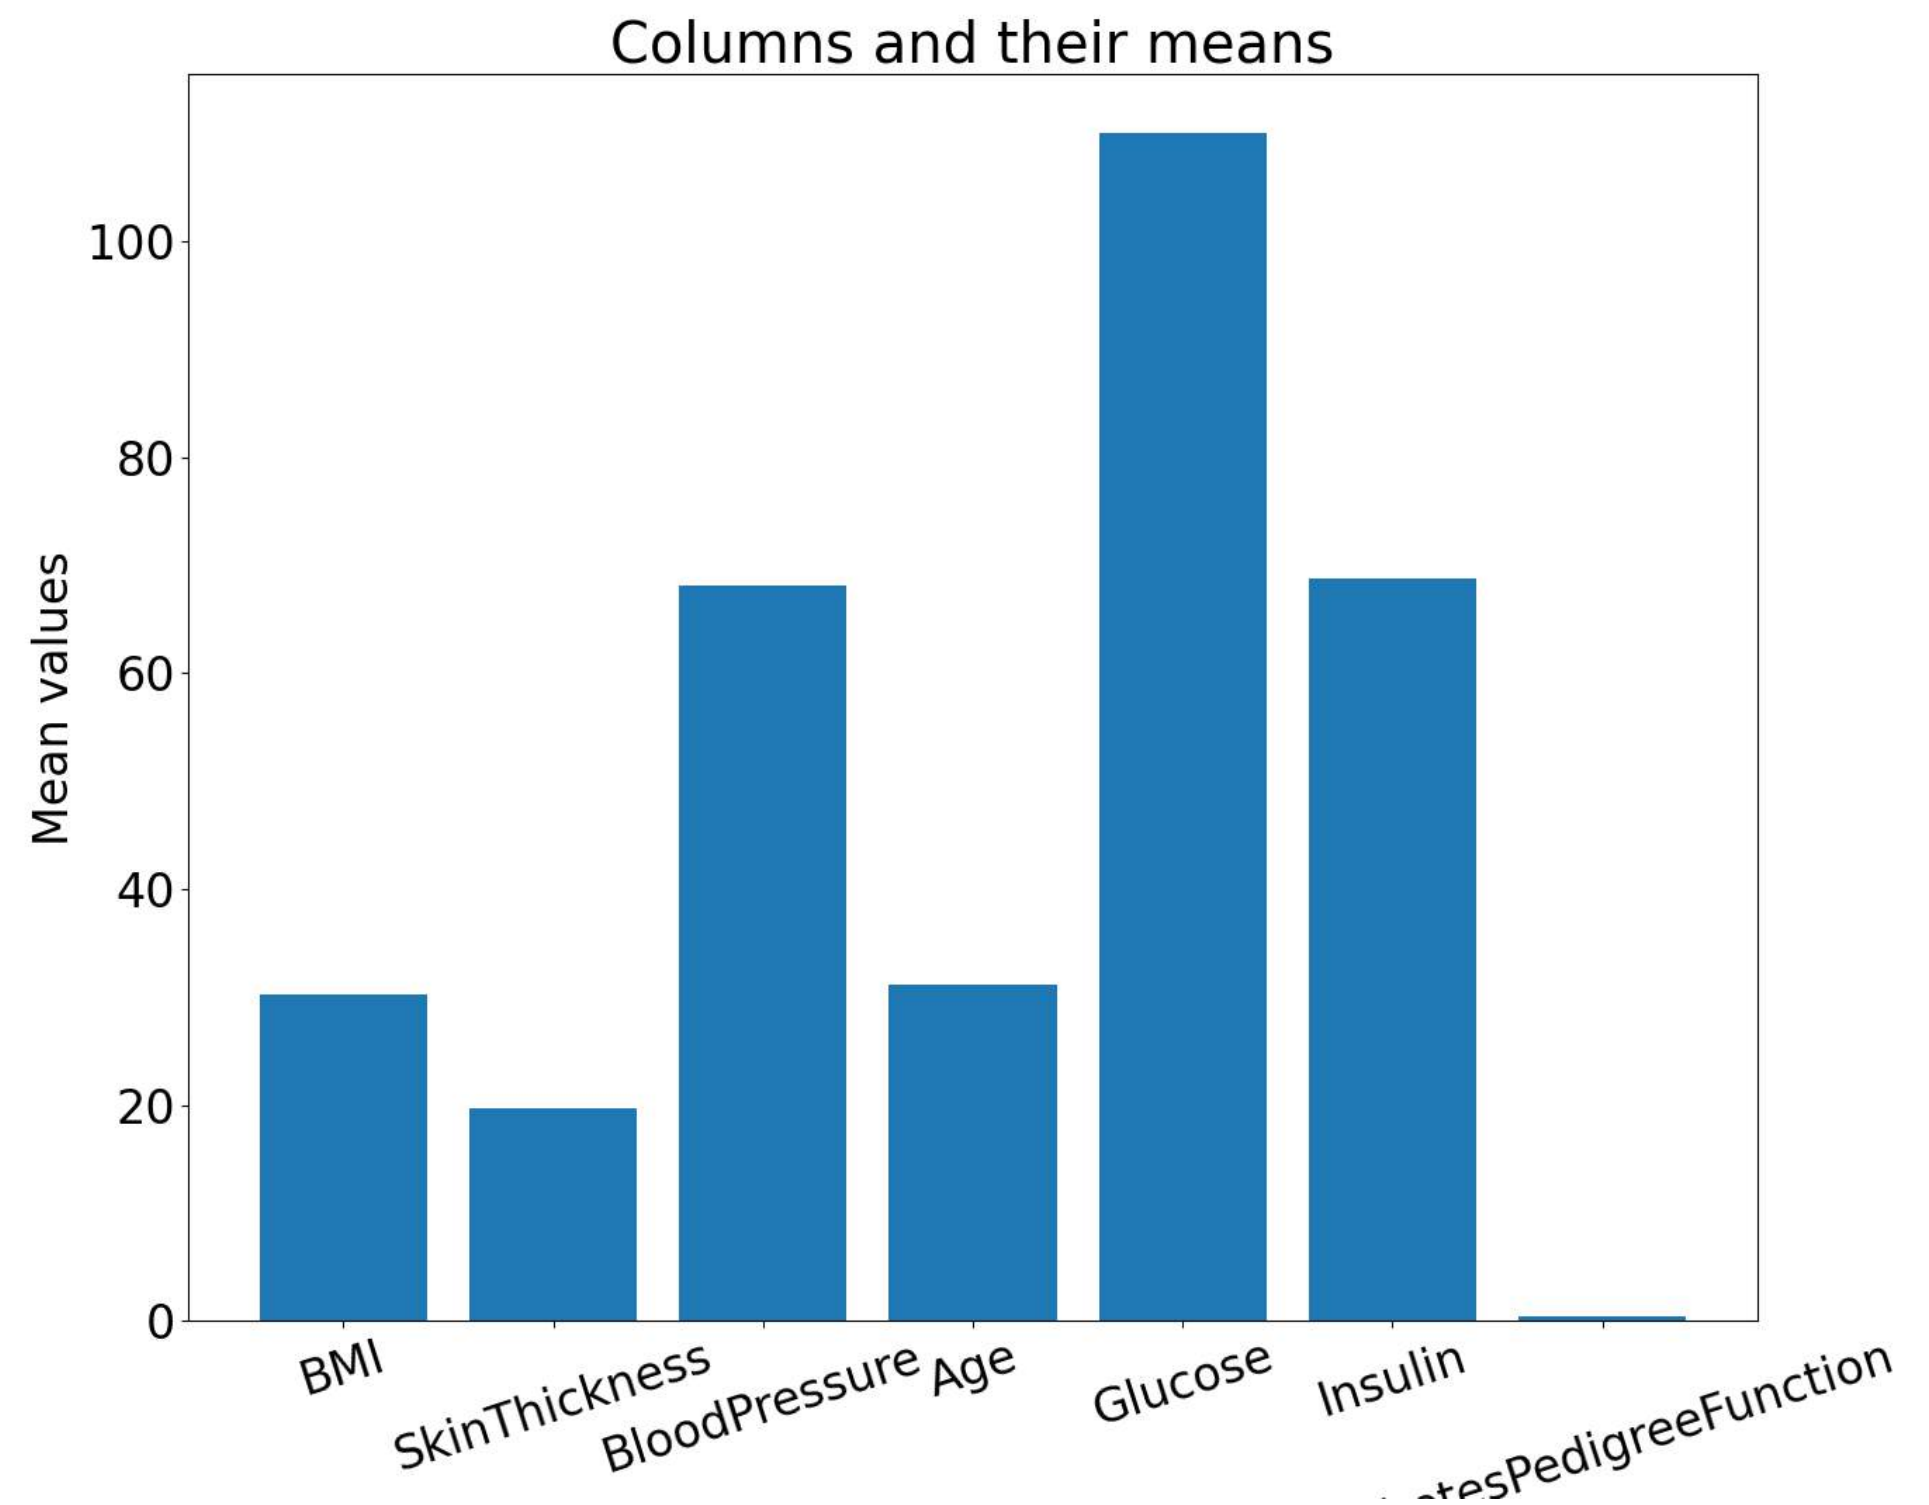
```
# Adjust figure size before plotting.  
plt.figure(figsize = (15, 12))
```

```
plt.bar(bar_positions, bar_heights)
```

```
plt.xticks(bar_positions,  
           bar_labels,  
           rotation = 18)
```

```
plt.ylabel('Mean values')
```

```
plt.title('Columns and their means') #  
<- add plot title
```



# Customize anything

- All possible style customizations are available in a `matplotlibrc` file
- **This sample** contains all of them, and any of those parameters can be passed to `rcParams` variable like we did earlier
- This sample contains a script of parameters and their default values
- Here's a part of that file with a sample of all parameters for modifying the style of the axes

```
## *****
## * AXES *
## *****
## Following are default face and edge colors, default tick sizes,
## default font sizes for tick labels, and so on. See
## https://matplotlib.org/api/axes_api.html#module-matplotlib.axes
#axes.facecolor:      white      # axes background color
#axes.edgecolor:      black      # axes edge color
#axes.linewidth:      0.8        # edge line width
#axes.grid:           False      # display grid or not
#axes.grid.axis:      both       # which axis the grid should apply to
#axes.grid.which:      major      # grid lines at {major, minor, both} ticks
#axes.titlelocation:  center     # alignment of the title: {left, right, center}
```

# Knowledge check



Link: [Click here to complete the knowledge check](#)

# Module completion checklist

Objective	Complete
Prepare data for visualization	✓
Create histograms, boxplots, and bar charts	✓



# Congratulations on completing this module!

You are now ready to try Tasks 1-13 in the Exercise for this topic

